

AgentNexLiFy — Product Plan

Author: Drafted with Claude (Cowork) on 27 May 2026 based on conversations with Aidan.
Status: Working draft. Decisions still required from owner are flagged inline.

1. The strategic shift, in one paragraph

We are consolidating the product around a single surface — **Agent OS** — and reclassifying every feature currently sold separately on the marketing site as a *template agent* (or supporting tool/data source) that owners access through natural language. The owner never opens a "SEO Audit" tab or a "Social Media" tool again. They talk to the orchestrator. The orchestrator picks the right template agent and runs it. The widget chatbot stays as the public customer-facing front door on the SMB's website, and every conversation it captures becomes data Agent OS agents can reference. Two surfaces, one shared data layer, one conversational entry point.

Sequencing decision (owner call, 27 May 2026): migrate every existing feature into the new model first ("best we can"), verify the migrated product works end-to-end, *then* curate templates into vertical packs. This preserves feature parity for existing paid customers, forces the template schema to handle the full breadth of the product, and defers vertical specialization until there are real templates to curate. The tradeoff: vertical defensibility comes later in the roadmap rather than sooner.

2. Where the product is today (factual audit)

2.1 What the marketing site currently promises

Four pricing tiers ranging \$0 to \$250/month, anchored on a list of nine features:

Feature today	Tier where it appears	Status in new model
AI chat widget	Free → Enterprise	Keep as public front door + data feed
Customer capture	Free → Enterprise	Implicit — lives in agent context
FAQ knowledge base	Free	Folded into agent context/KB

Feature today	Tier where it appears	Status in new model
Hosted business page	Free → Enterprise	Open decision — deprecate, or keep as a background asset agents can update
Appointment booking	Growth+	→ Booking template agent (needs real Google Calendar integration)
Auto follow-up email & SMS	Growth+	→ Lead Nurture template + Follow-up trigger
Automation sequences (2 / 6 / unlimited)	Growth+	→ Trigger configurations on templates
AI content writer	Growth+	→ Content template agent
Basic SEO audit	Growth+	→ SEO Audit template agent (needs actual crawler backend)
Full SEO suite & keyword tracking	Pro+	→ Same agent, deeper data sources at higher tiers
Social media content & scheduling	Pro+	→ Social Post template agent (drafts now, scheduling later)
Email & SMS marketing campaigns	Pro+	→ Campaign template agent + send action
Review request automation	Pro+	→ Review Request template + trigger
Pipeline automation	Pro+	→ Triggers wired into agent context, not a separate UI
AI marketing content generator	Pro+	→ Folded into Campaign + Content templates
Analytics & reporting	All tiers, varies	→ "Weekly Briefing" template agent that produces summaries on a schedule

Feature today	Tier where it appears	Status in new model
AI appointment booking agent	Enterprise	→ Baseline capability of Booking template at higher tiers (autonomous booking)
Team accounts & roles	Enterprise	→ Stays as platform feature
AI visibility tracking (GEO score)	Enterprise	→ AI Visibility template agent
Priority onboarding	Enterprise	→ Stays as service
White-label & custom CSS	Enterprise	→ Stays as platform feature
Webhook integrations	Enterprise	→ Stays as platform feature (also unlocks agent actions to external tools)

Implication: roughly 15 of the 22 paid features map cleanly to a template agent. A handful (team accounts, webhooks, white-label, priority onboarding) stay as platform-level features rather than agents. Two (hosted business page, dashboard analytics) are decisions: deprecate, or rework so agents drive them.

2.2 What's actually in the product right now (per my testing)

- Marketing site: hero, How It Works (3 steps), Features (9 cards), Pricing (4 tiers), FAQ, Demo CTA. Two typos in the H1.
- Signup → onboarding wizard (38% complete on test account) → dashboard. Industry pre-configures the dashboard to that vertical (e.g. "Configured for Auto Shop").
- Sidebar today: Dashboard, Analytics, Agent Control, Agent OS, CRM (Clients / Pipeline / Smart Lists / Client Portal), Communications, Marketing, Operations.
- Agent OS: orchestrator + 5 worker agents (Booking, Customer Question, Lead Nurture, Campaign, Generalist). All drafts-only, with an Approve/Reject side panel and a reasoning trace.
- Widget chatbot on every page (including the signup page, where it competes with the form — see QA report).

2.3 What the QA report flagged as material issues

(Carry these forward into the plan — they're sequencing inputs.)

- **Signup data isn't piped to agent prompts.** Drafts come back with [Shop Name], [Your Name], [Phone] placeholders.

- **Reasoning trace is partially theatrical.** Workers claim to "Load business profile / knowledge base / recent chats" even when those are empty.
- **Customer Question agent can produce owner-facing text inside a customer-labelled draft.** Real risk if any auto-send ships.
- **Out-of-credits state is invisible** to the owner; produces silent empty drafts.
- **Generalist agent produced empty output** for both analytical and simple-writing prompts during my session (some of this was credit-related).

These get prioritized into Phase 1 of the roadmap below.

2.4 Feature categorization (four buckets, not just "templates")

Not every feature should become a template. The migration produces four kinds of artifacts:

Bucket	What it is	Examples from current product
Template agents	Generative units of work an agent does end-to-end	Booking drafts, Lead Nurture sequences, Customer Q&A replies, Content Writer, Campaign emails, Review Requests, Social Posts
Tools / capabilities	Things templates <i>call</i> — not user-facing on their own	Google Calendar integration, email send (Gmail OAuth), SMS send (Twilio), Stripe payment links, SEO crawler service (future)
Data sources	Information that flows into the shared context for any template to reference	Widget conversation history, business profile, lead/pipeline state, agent run history
Platform features	Account-level functionality with no agent involvement	Team accounts & roles, webhooks, white-label branding, billing

This categorization keeps the template registry clean (a template is a unit of agent capability) and makes it obvious when a "feature" actually needs a service backend (e.g. SEO Audit) before any template can be useful.


2.5 The "best we can" features


Three features genuinely don't fit the template paradigm cleanly. Plan for friction during migration:



- **SEO Audit / SEO Suite.** Needs a real crawler. Short-term migration: a template that returns generic best-practice recommendations based on the owner's URL. Real version requires building or buying a crawl service. Be honest in marketing during migration — call it "SEO recommendations," not "SEO audit."
- **Analytics & Reporting.** Not a generative task — it's data aggregation and visualization. Replace with a **Weekly Briefing template** that reads agent run history + widget data + pipeline state and produces a written summary on Mondays. The numbers dashboard goes away; the insight stays.
- **AI Visibility tracking (GEO score).** Needs an actual data source (querying LLM-based search results). Stub during migration; real version requires real infrastructure or a partnership.

The **Hosted Business Page** is a fourth edge case. It's an artifact (a hosted webpage), not a generative output. Options: (a) deprecate cleanly; (b) keep the page and build a **Page Editor template** that updates its copy via conversation. *Owner recommendation in §7 is deprecate, but the call is open.*

3. Target architecture







3.1 The five architectural pieces

1. Orchestrator (Haiku-powered).

- One job: classify the owner's natural-language request to one of N templates (or no-match).

- Cheap, fast, fine-tunable on real owner data over time.
- Surfaces the routing decision to the owner ("I'm picking the booking template — sound right? [pick another]").
- When no template matches → captures as wishlist (see 3.3) and offers nearest-match fallback.

2. Template Agent Registry.

- Each template has a strict schema: `name`, `description`, `when_to_use`, `param_schema`, `prompt_scaffolding`, `trigger_config`, `permission_scope`, `expected_outputs`, `example_runs`.
- Templates are versioned. Owners' saved configurations pin a version.
- Templates declare what data they need to read from the shared context and what they're allowed to write/send.

3. Vertical Packs.

- Curated bundles of templates appropriate for one industry.
- Owner's signup industry pre-selects the pack; owner can switch later.
- This is where the marketing story changes: "The Auto Shop Pack — 25 agents that handle everything from booking to reviews to seasonal campaigns" beats "AI for SMBs" cold.
- Vertical pack composition becomes a competitive moat over time.

4. Shared Context / Data Layer.

- Single source of truth for everything an agent might reference.
- **Critically: widget conversations flow in here automatically.** When the widget captures a lead chat about "weak hybrid battery on a 2018 Prius," that becomes a record the Lead Nurture template can act on tomorrow.
- Also stores: business profile (the signup data we are currently failing to use), pipeline state, agent run history.
- Reasoning trace UI reads from here and shows the owner *what data was actually used* — not a theatrical "Loaded" indicator.

5. Triggers.

- Templates can be invoked manually (owner asks orchestrator), on a schedule (every Monday morning), or on an event (new widget conversation, new lead, lead stale for >3 days, etc.).
- Trigger config is part of every template's schema from day one, even if Phase 1 ships only manual.

- This is the single architectural decision that turns Agent OS from a chat assistant into a real platform. Plan it now even if you don't ship it for two quarters.

3.2 The widget as a first-class data source

Concrete contract:

- Widget conversation finishes → ingestion job summarises it (intent, contact info, vehicle/product details, outcome) into the shared context.
- Every widget conversation becomes addressable by agents: "the widget thread from Mike on Tuesday."
- Triggered templates can subscribe to widget events: a "new widget lead" event can fire a Triage template that classifies and slots the lead into the pipeline; a "widget conversation closed without booking" event can fire a Follow-up template 24 hours later.
- Owner can ask the orchestrator: "what came in through the widget yesterday?" — and the orchestrator can answer directly without involving a worker.

3.3 The wishlist mechanism

When the orchestrator's classifier returns low confidence on all templates:

1. **Acknowledge in chat** — "I don't have a template that fits this exactly. The closest match would be [X] — want me to try, or would you rather I save this request and we build a proper template for it?"
2. **Capture wishlist entry** — record the verbatim request, the matched-templates-considered list, the owner's vertical, and frequency (how often does this kind of request come in?).
3. **Internal prioritization signal** — wishlist requests aggregate into a product-side dashboard you can sort by vertical and frequency. This becomes your roadmap for new templates *driven by real customer demand*, not internal guessing.
4. **Notify the owner when a matching template ships** — "you asked for X three weeks ago; here's the new template that handles it."

This is a single feature that does multiple jobs: handles the long tail gracefully, gives owners a sense of agency, gives you free product research, and creates a delighter when the template later appears.

4. Pricing reframe

Timing note (per owner sequencing decision): the pricing rewrite happens *after* migration + verification + vertical pack curation. Until then, today's tier structure

stays — existing customers continue to pay for the features they were sold; those features now live inside Agent OS instead of separate dashboard pages.

The current tier model rewards *more features*. The new model needs to reward *more agent usage and more autonomy*. Concrete proposal (numbers tunable, structure is the point):

Tier	Price	Agents	Autonomy	Triggers	Data
Free	\$0	Vertical pack (drafts-only)	Drafts only — owner must copy/paste to send	Manual only	7-day widget history
Growth	\$99	Vertical pack (drafts + real actions)	Real send for email + calendar booking	Manual + scheduled	90-day history, 200 agent runs/mo
Pro	\$150	All vertical packs	All actions enabled	Manual + scheduled + event-based	1-year history, 500 runs/mo, custom template tuning
Enterprise	\$250	All packs + custom templates	Unlimited actions	Everything	Unlimited history & runs, API access, team accounts, white-label

The principle: owners on every tier get the *same surface* (Agent OS). What changes is how much it can do for them. Free is provably useful (drafts they can copy/paste). Paid tiers turn on autonomy, integrations, and frequency.

Important migration concern: existing Growth/Pro/Enterprise customers were sold on specific features (SEO suite, social scheduling, etc.). They need a transition story that preserves their access to those capabilities via agents. Concretely:

- Email all paid customers ~4 weeks before launch.
- Tell them: same price, same capability, much easier UX. Show 1 GIF of "ask the agent to run an SEO audit" vs the old way.

- Keep the old UI live for 60 days post-launch behind a "classic view" toggle so they can ramp at their own pace.
-

5. Marketing site reframe

Timing note: the full site rewrite waits until vertical packs ship. During migration the existing site stays live, *except* for the Phase 1 fixes (typos, hide widget on /signup). The hero copy and Features grid can be revised once Agent OS-through-chat is the primary experience.

The site today undermines the new positioning. The "Features" grid (9 cards) implies a multi-tool product. New site structure (target):

- **Hero:** "Run your business by talking to your AI" (or similar — single-promise framing).
 - **The demo:** one short interaction showing the owner asking the orchestrator for something, the orchestrator picking a template, the agent running, the owner approving. This is the product. Don't show a feature grid.
 - **Vertical landing pages:** [/auto-shops](#), [/dental](#), [/hvac](#), [/salons](#). Each shows the pack's templates with example owner asks ("Hey, send a follow-up to everyone we quoted last week who didn't book"). This is where conversion will happen.
 - **Pricing:** simpler. Three or four tiers, anchored on autonomy + usage, not feature counts.
 - **FAQ:** keep "what if AI makes a mistake?" — that's a strong existing answer. Update the other FAQ entries to match the new model.
 - **Trust signals:** at least one real customer quote, with a number. ("12 extra appointments this month at Sunset Mobile Detailing.") Replace the bot mockup with a real screenshot of Agent OS.
 - **Fix the two H1 typos** before any of the above ships.
-

6. Phased roadmap (revised for migration-first sequencing)

The shape: substrate fixes → full feature migration → verification → real actions → vertical curation → pricing & marketing reframe. Vertical packs come *after* every existing feature lives in Agent OS, which means the first six months are about parity, and competitive defensibility comes later in the year.

Phase 1 — Substrate fixes (weeks 1–6)

Non-negotiable foundation. None of the migration work below will pay off if these aren't done first. Don't ship new features in this phase; fix what's broken.

1. **Wire signup data into agent prompts.** Business name, owner name, industry, city, phone, website should be in every worker's prompt. Single highest-leverage fix in the product — eliminates [Shop Name]-style placeholders across every template you'll build later. (1 week)
2. **Honest reasoning trace.** Replace "Loading business profile (Hours, services, timezone)" with the actual data that was loaded, shown inline. If nothing was loaded, say so. (3–5 days)
3. **Surface credit/usage state.** Account banner when credits low; explicit "service temporarily unavailable" in-draft message when LLM is unavailable. Never produce silent empty drafts. (3 days)
4. **Fix the worst signup bugs.** Two H1 typos, stale "Passwords do not match" error, phone validation, hide the widget on /signup. (2 days)
5. **Swap orchestrator to Haiku.** Classification/routing on Haiku; Sonnet only for final drafts. Measure cost-per-task before and after. (1 week)
6. **Build the wishlist capture mechanism.** Simple version: when orchestrator low-confidence, prompt the owner, log to a wishlist table, internal-only dashboard. Email notifications later. (1 week)
7. **Lock the template schema.** Even though templates aren't being built yet, define the schema now so every migration in Phase 2 conforms. Fields: name, description, when_to_use, param_schema, prompt_scaffolding, trigger_config (default manual), permission_scope, tool_dependencies, expected_outputs, example_runs. (3 days)

Exit criteria for Phase 1: the existing 5-agent product produces drafts that include real business data, the reasoning trace honestly reflects what was loaded, the credit-out failure mode is visible to the owner, and the template schema is documented. No new templates yet.

Phase 2 — Full feature migration to templates + tools (weeks 7–18)

The big one. Twelve weeks to migrate every existing paid feature into the new model. Build in approximate order of customer importance — most-used features first, riskiest "best we can" features last.

2a. Migrate the existing 5 worker agents into the template schema (weeks 7–8). Booking, Customer Question, Lead Nurture, Campaign, Generalist. No user-facing change. This is a refactor that proves the schema works. While doing it, fix the bugs from the QA report: SMS markdown stripping, the Customer-Question-asking-owner failure mode, the Booking logical contradiction.

2b. Build the parameter-extraction layer (week 9). Orchestrator extracts structured params from the natural-language request and hands them as typed inputs to the template. This is what separates a template platform from "prompts with a router."

2c. Migrate the "clean template" features (weeks 10–13). Each is a new template in the registry. All start as drafts-only — real actions land in Phase 4.

- **Content Writer** → Content template (blog posts, social copy, marketing content).
- **Review Request** → Review Request template + (later) post-service trigger.
- **Social Post** → Social Post template.
- **Email/SMS Campaign** → already a template (Campaign); validate it covers the existing feature scope.
- **Quote Follow-up / Lead Nurture** → already a template; validate timeline handling (fix the "Day 5 vs. 'a little over a week ago'" issue I flagged in QA).
- **Auto follow-up email & SMS** → not new templates per se; these become *triggers* on existing templates. Trigger plumbing happens in Phase 4, but tag now which templates need which triggers.
- **Pipeline automation** → not a template. Pipeline state becomes structured data in the shared context; agents reference it. Owner-facing pipeline UI either deprecates or becomes a read-only view of agent-driven state.

2d. Build the tools layer (weeks 11–14, in parallel). These are *not* templates — they're capabilities templates call. Build them as services with clear API contracts that any future template can use.

- **Google Calendar tool.** Read-only first (templates can see availability), then write (templates can create appointments).
- **Email-send tool.** Gmail OAuth first; SMTP as fallback. Send-on-behalf with audit logging.
- **SMS tool.** Twilio integration. Owner provides their own Twilio number to start.
- **Widget-conversation-search tool.** Templates can query the captured widget data ("the lead from Mike about hybrids on Tuesday").

2e. Migrate the "best we can" features (weeks 14–17). These have known limitations — ship them with clear marketing language about what they actually do.

- **SEO Recommendations template.** Takes a URL, returns generic best-practice recommendations + checks meta tags / heading structure / mobile viewport. *Real crawler is a Phase 5+ project.*
- **Weekly Briefing template** (replaces Analytics dashboard). Reads agent run history + widget data + pipeline state, writes a Monday-morning summary. Default trigger: every Monday 7am owner time.
- **AI Visibility template.** Stub for now; returns "this feature is in early access — we're tracking how AI search engines describe your business; full results coming soon" plus any manual signals we can grab. *Real version is a partnership or build-out decision.*

- **Page Editor template** (if Hosted Business Page stays). Owner says "update my page to mention the AC special"; template proposes edits to a specific page section; owner approves. *Otherwise, deprecate Hosted Business Page in this phase.*

2f. Platform features (weeks 14–18, in parallel). No agent involvement; keep these where they are.

- Team accounts, webhooks, white-label, billing — confirm they continue to work as Agent OS becomes the primary surface.

Exit criteria for Phase 2: every feature on the current pricing page has a corresponding template or tool. Owner can do everything they could before, through Agent OS. Old dashboard pages still exist but are de-emphasized.

Phase 3 — Verification & hardening (weeks 19–22)

Don't build new things in this phase. Make sure what you built actually works.

1. **End-to-end test every template.** One realistic owner ask per template, full run, validated output. Catalogue gaps.
2. **Routing accuracy audit.** Sample 200 real owner asks (from beta or staging traffic). Did the orchestrator pick the right template? Where it didn't, retrain or refine.
3. **Cost audit.** Per-template cost-per-run measured. Anything that costs more than 5× the median needs investigation.
4. **Failure-mode audit.** Re-test every issue from the QA report against migrated templates. Confirm the substrate fixes hold under load.
5. **Internal dogfood.** Run your own business operations through Agent OS for two weeks. Anything you don't trust enough to actually use isn't ready.
6. **Migrate one friendly paid customer.** Walk them through the new experience. Capture their reaction, ship fixes from it.

Exit criteria for Phase 3: internal team uses Agent OS daily. Routing accuracy $\geq 90\%$ on realistic test set. No P0/P1 bugs open.

Phase 4 — Real actions (weeks 23–30)

Take the product from "drafts" to "agents that actually do things," with safety first.

1. **Permission scopes per template.** Owner-configurable: "this template's drafts require approval / auto-send"; "can send to existing customers only"; "spending cap \$X/month"; "rate cap N sends/day." *Build this before turning anything on.*
2. **Action audit log.** Every action by every agent: timestamp, template version, params, output, outcome. One-click view, one-click revert where possible.

3. **Calendar bookings actually book.** The Booking template, when approved, writes to the connected Calendar. This is the single highest-impact change in the entire roadmap from the owner's perspective.
4. **Email send goes live.** Approved drafts can actually send via Gmail OAuth. Starts with a hard cap (e.g. 20 sends/day per account) and expands as confidence builds.
5. **SMS send goes live.** Same pattern via Twilio.
6. **Scheduled triggers.** Templates can run on schedules (every Monday 7am, every Tuesday 10am, etc.). Weekly Briefing template moves from manual to scheduled by default.

Exit criteria for Phase 4: owner can configure an agent to actually send an email after approving the draft, and the product remembers/audits what was sent.

Phase 5 — Vertical pack curation (weeks 31–38)

Now the strategic work — the templates and tools exist; turn them into vertical packs.

1. **Define the pack format.** A pack is: a labelled bundle of templates, vertical-specific tone/style overrides on shared templates, vertical-specific example runs in onboarding, vertical-specific suggested triggers. Schema-wise: a thin layer on top of the existing template registry.
2. **Curate the Auto Shop Pack.** Pick the relevant templates from the registry (~15–20 of the ~25 you'll have). Tune copy/tone for auto-shop language. Suggest auto-shop-appropriate triggers (oil-change reminders, seasonal campaigns). Build 3 example run-throughs for onboarding.
3. **Curate the HVAC Pack** — same exercise.
4. **Curate the Dental Pack** — same exercise.
5. **Pre-select pack by industry at signup.** Owner picks "Auto Shop" → gets the Auto Shop Pack pre-loaded with relevant templates active and tuned.
6. **Event-based triggers** (new widget lead → triage; lead stale 3 days → follow-up). This is where the product starts to feel proactive rather than reactive.

Exit criteria for Phase 5: a new owner who signs up as "Auto Shop" sees a curated, vertical-tuned agent suite from minute one.

Phase 6 — Pricing & marketing reframe (weeks 39–44)

Now that there's a clearly differentiated product (vertical packs, real actions, autonomous triggers), retire the feature-list pricing and feature-grid marketing.

1. **New pricing model goes live.** Tiers by autonomy and usage, not feature count. Existing customers grandfathered to comparable tiers automatically.

2. **Marketing site rewrite.** Hero promises the single thing the product now does. Vertical landing pages ([/auto-shops](#), [/hvac](#), [/dental](#)). Replace the feature grid with example owner asks per vertical.
3. **One real-customer trust signal per vertical.** Number-backed case study. Three packs, three real customers cited.

Exit criteria for Phase 6: marketing matches product reality. New visitors see a focused single-promise pitch.

Phase 7+ — Customization, marketplace, expansion (weeks 45+)

1. Owners can save custom-tuned versions of templates ("my version of the Lead Nurture template").
 2. Eventually, owner-created templates from scratch (with safe-defaults from the schema).
 3. Vertical expert invite-only program — contribute templates to packs you know.
 4. API access for Enterprise customers who want to embed the orchestrator elsewhere.
 5. Real SEO crawler infrastructure (or partnership).
 6. Real AI visibility tracking (or partnership).
 7. Additional vertical packs prioritized by wishlist data.
-

7. Risks and open decisions

Risks to plan around

1. **Customer migration risk.** Existing paid customers were sold on specific features. Until Phase 6, their tier names and feature promises remain. After Phase 6's pricing reframe they need a transition story that preserves their access. Test on 5 friendly customers before broad rollout.
2. **Routing accuracy = product quality.** If the Haiku-powered orchestrator misroutes 1 in 10 requests, owners will lose trust. Plan for: showing routing decisions to owners, instrumenting every routing decision, fine-tuning a small classifier on real owner data after 1000+ routing decisions are logged.
3. **Cost-per-task at scale.** Free tier is dangerous if it produces unlimited Sonnet drafts. Hard caps per tier (runs/month) are non-negotiable. Track cost-per-active-user weekly. *Especially relevant in Phase 2 when 15+ templates ship and per-template costs vary widely.*
4. **The "no template fits" UX is the long tail.** Owners will ask for things templates don't cover. Wishlist handles this, but Phase 1 needs a smooth nearest-match fallback so it doesn't feel broken.
5. **Real-action safety.** The moment agents start sending things, you need a safety architecture: scoped permissions per template, spending/rate caps, sender allow-lists,

content checks, kill switch. Build the safety architecture *before* turning on real send (Phase 4 sequences this correctly).

6. **Vertical pack maintenance.** Each pack is 15–25 templates. Three packs = up to 75 templates. Versioning, regression testing, keeping vertical-specific copy fresh — this is real engineering load that grows linearly with packs.
7. **Migration-first specific: feature regression risk.** Owners on paid tiers expect every feature they were sold to keep working through migration. If a migrated template is materially worse than its predecessor for the first 2–4 weeks of Phase 3, churn risk spikes. Mitigation: hard verification gate (Phase 3) before retiring any old UI, and keep old dashboard pages reachable via "classic view" toggle through Phase 6.
8. **Migration-first specific: horizontal-stall risk.** Twelve weeks of feature-migration work without vertical differentiation means competitors with better distribution (Podium, ServiceTitan) keep their lead. Mitigation: even during Phase 2, keep one piece of vertical work alive in parallel — at minimum the example runs / onboarding copy for the auto-shop industry that's already pre-configured at signup.
9. **Migration-first specific: "best we can" credibility risk.** Stub SEO Audit and Visibility Tracking templates risk owner disappointment if they expect the full feature. Mitigation: be honest in product copy ("SEO recommendations — full audit coming Q3") and don't sell upgrade tiers on these specific features during migration.

Decisions still needed from the owner

1. **Hosted Business Page** — deprecate cleanly, or build a Page Editor template that maintains it? *Recommend: deprecate. The marketing value is small; the maintenance cost is real. Decide before Phase 2e starts.*
 2. **First vertical pack to curate** (Phase 5). *Recommend: Auto Shops* (already wired into signup; QA scenarios are anchored on this vertical). Alternatives: HVAC (huge year-round market) or dental (higher willingness-to-pay).
 3. **"Classic view" toggle during migration?** Strongly recommend yes — keeps old dashboard pages reachable through Phase 6 in case a migrated template falls short. Decision needed: hard cutover, or 60-day toggle, or longer?
 4. **Drafts-only on Free tier indefinitely, or with a usage cap?** *Recommend: drafts-only indefinitely. It's the safety story and the upgrade hook.*
 5. **Real send architecture (Phase 4) — Gmail OAuth first, or owner-provided SMTP?** *Recommend: Gmail OAuth + Outlook OAuth covers ~80% of SMB email. SMTP later.*
 6. **Naming.** "Agent OS" is the right product name, but "AgentNexLiFy" as the company name is a mouthful. Worth a marketing-side conversation. (*Out of scope for this plan, but flagging.*)
-

8. Success metrics to track from week one

- **Free → paid conversion rate.** Today's tiered-feature model will look very different from a usage/autonomy model.
 - **Routing accuracy.** % of orchestrator routings that the owner accepts without re-routing.
 - **Wishlist volume per vertical.** Top 10 wishlist items per pack drive template-building priorities.
 - **Time-to-first-draft from signup.** Should drop to <5 minutes; today it's gated on the 8-step onboarding wizard.
 - **Cost per active agent run.** Watched weekly. If this trends up while pricing stays flat, free tier becomes a problem.
 - **Owner-approved actions per week** (once real actions ship). This is the long-run leading indicator that the product is doing real work.
-

9. The single sentence to test this plan against

"A non-technical small business owner can sign up, pick their industry, and within five minutes ask their AI in plain English to do something useful — and the AI either does it, drafts it for approval, or tells them honestly it doesn't have a template for that yet and saves the request for next time."

If a feature, decision, or piece of polish doesn't move us toward that sentence, it doesn't make Phase 1 or 2. Everything else is later.

— end of plan —